

1.0 Importing Libraries

- Setting various of jupyter notebook for better readability

```
In [1]: import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
from matplotlib.pyplot import figure, xticks
import seaborn as sns

from sklearn.preprocessing import MinMaxScaler, StandardScaler, scale, PolynomialFeatures
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, RandomizedSearchCV, cross_validate, KFold
from sklearn.linear_model import Ridge, RidgeCV, Lasso, LassoCV, LinearRegression, ElasticNet
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier, export_graphviz
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, BaggingRegressor, GradientBoostingRegressor
from sklearn.metrics import r2_score, precision_score, recall_score, mean_squared_error, mean_absolute_error
from sklearn.feature_selection import RFE
from sklearn import preprocessing, utils

import xgboost as xgb
from xgboost import XGBRegressor
import lightgbm as lgbm
from lightgbm import LGBMRegressor

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

from IPython.display import Image
from six import StringIO
from io import StringIO
import pydotplus
import graphviz

pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)
```

2.0 Importing Dataframe and evaluating basic properties

```
In [2]: base=pd.read_csv('CTGAN Generated data.csv')
base=base.rename(columns=lambda x: x.strip())
```

```
In [3]: base.head()
```

Out[3]:

	Feed N Plus 2A content	Reactor WAIT	H2 to HC	Reactor 1 Inlet Temp	Reactor 2 Inlet Temp	Reactor 3 Inlet Temp	Reactor 4 Inlet Temp	Reactor 1 Delta T	Reactor 2 Delta T	Reactor 3 Delta T	Reactor 4 Delta T	Reactor 1 Delta P	Reactor 2 Delta P	Reactor 3 Delta P	Reactor 4 Delta P	Seperator Pressure	Seperator Temperature	Recycle gas purity	Net gas Hydrogen Purity	Coke on Spent Catalyst	Chloride Injection rate	Total Paraffins in feed	Total Naphthenes in feed	Total Aromatics in feed	Total olefin I Fee
0	45.46	1004.48	3.28	998.41	1000.46	1006.41	1011.81	166.26	108.39	72.02	39.17	1.17	2.73	2.77	3.17	30.83	100.62	78.67	88.24	3.82	2.59	64.91	24.60	10.43	0.0
1	44.22	1004.54	3.17	999.00	999.57	1003.23	1012.77	166.37	105.33	68.00	38.45	1.17	2.80	2.78	3.22	30.32	103.02	78.67	88.24	3.56	2.64	65.89	23.58	10.32	0.2
2	44.22	1004.54	3.17	996.97	999.72	1004.46	1012.81	165.14	107.33	69.92	39.13	1.16	2.74	2.77	3.25	30.27	103.55	78.67	88.24	3.56	2.58	65.89	23.58	10.32	0.2
3	45.49	1002.61	3.46	998.34	997.32	1001.66	1011.66	172.18	106.72	70.46	40.10	1.25	2.72	2.76	3.18	32.56	103.20	77.88	88.40	3.82	2.60	63.35	27.07	9.21	0.3
4	45.46	1004.34	3.31	998.17	999.71	1004.88	1013.84	165.81	107.55	69.72	39.09	1.16	2.77	2.76	3.09	30.39	103.05	78.67	88.24	3.82	2.60	64.91	24.60	10.43	0.0

```
In [4]: base.shape
```

Out[4]: (5962, 30)

3.0 Data Cleaning

```
In [5]: base.describe()
```

Out[5]:

	Feed N Plus 2A content	Reactor WAIT	H2 to HC	Reactor 1 Inlet Temp	Reactor 2 Inlet Temp	Reactor 3 Inlet Temp	Reactor 4 Inlet Temp	Reactor 1 Delta T	Reactor 2 Delta T	Reactor 3 Delta T	Reactor 4 Delta T	Reactor 1 Delta P	Reactor 2 Delta P	Reactor 3 Delta P	Reactor 4 Delta P	Seperator Pressure	Seperator Temperature	Recycle gas purity	Net gas Hydrogen Purity	Coke on Spent Catalyst	Chloride Injection rate	Total Paraffins in feed	Total Naphthenes in feed	Total Aromatics in feed	Total olefin I Fee
count	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000	5962.000000
mean	47.972068	982.563648	3.814648	975.295392	982.172801	983.606872	986.789906	161.030943	106.133452	71.095792	41.762420	0.992466	2.587736	2.640976	2.805659	33.360704	105.697828	81.232	88.24	3.56	2.64	65.89	23.58	10.32	0.2
std	2.046191	14.284562	0.335320	13.532721	15.434145	13.933120	15.674754	7.303361	6.186524	3.766175	4.140208	0.341575	0.091843	0.132266	0.240884	1.310147	3.924914	9.575	88.24	3.56	2.58	65.89	23.58	10.32	0.2
min	40.380000	963.040000	3.100000	945.740000	951.400000	960.060000	964.350000	134.380000	86.240000	52.290000	22.690000	0.450000	2.330000	2.150000	2.130000	28.650000	92.960000	0.860	88.24	3.56	2.58	65.89	23.58	10.32	0.2
25%	47.000000	968.530000	3.560000	964.770000	966.470000	970.472500	971.550000	156.232500	101.560000	69.530000	39.670000	0.650000	2.520000	2.530000	2.640000	32.510000	103.290000	80.880	88.24	3.56	2.58	65.89	23.58	10.32	0.2
50%	48.150000	980.890000	3.770000	970.845000	981.850000	981.610000	986.385000	161.160000	106.120000	71.350000	42.580000	1.100000	2.570000	2.650000	2.840000	33.490000	105.805000	82.550	88.24	3.56	2.58	65.89	23.58	10.32	0.2
75%	49.380000	991.927500	4.030000	984.447500	994.270000	995.450000	996.162500	166.407500	110.450000	73.137500	44.670000	1.310000	2.670000	2.750000	2.980000	34.300000	108.307500	83.450	88.24	3.56	2.58	65.89	23.58	10.32	0.2
max	59.890000	1009.490000	5.660000	1004.350000	1015.790000	1009.320000	1015.620000	176.950000	128.670000	83.260000	51.740000	1.880000	2.820000	2.920000	3.280000	36.500000	119.300000	86.820	88.24	3.56	2.58	65.89	23.58	10.32	0.2

```
In [6]: ### Removing outliers using Q1-1.5IQR & Q3+1.5IQR
def mod_outlier(df):

    df1 = df.copy()
    df1 = df1._get_numeric_data()

    df2 = df.copy()

    q1 = df1.quantile(0.1)
    q3 = df1.quantile(0.9)

    iqr = q3 - q1

    lower_bound = q1 -(1.5 * iqr)
    upper_bound = q3 +(1.5 * iqr)

    for col in df.columns:
        for i in range(0,len(df1[col])):
            if df1[col][i] < lower_bound[col]:
                df1[col][i] = np.nan

            if df1[col][i] > upper_bound[col]:
                df1[col][i] = np.nan

    for col in df.columns:
        df2[col] = df1[col]

    return(df2)

df_outlier_removed = mod_outlier(base) #Passing the base dataframe through the mod_outlier function
print(df_outlier_removed.isna().sum())
df_outlier_removed=df_outlier_removed.dropna() ## Dropping blank rows
```

```
Feed N Plus 2A content      18
Reactor WAIT               0
H2 to HC                   13
Reactor 1 Inlet Temp        0
Reactor 2 Inlet Temp        0
Reactor 3 Inlet Temp        0
Reactor 4 Inlet Temp        0
Reactor 1 Delta T           0
Reactor 2 Delta T           0
Reactor 3 Delta T           2
Reactor 4 Delta T           0
Reactor 1 Delta P           0
Reactor 2 Delta P           0
Reactor 3 Delta P           0
Reactor 4 Delta P           0
Seperator Pressure          0
Seperator Temperature       0
Recycle gas purity          80
Net gas Hydrogen Purity     0
Coke on Spent Catalyst      0
Chloride Injection rate     2
Total Paraffins in feed     21
Total Naphthenes in feed    18
Total Aromatics in feed     0
Total olefins in Feed       142
Reactor LHSV                22
Feed IBP                    0
50% IBP                     69
WABT                        0
Plant C5PlusYield           0
dtype: int64
```

4.0 Identifying Features/Independent variables for building models

4.1 Feature Selection Based on RFE using DecisionTreeRegressor and LinearRegression

```
In [7]: # RFE using DecisionTree Regressor

def rfe_DecissionTreeRegressor(df, target_col, k):
    # Create a copy of the original dataframe
    df_new = df.copy()

    # Get the target column values
    y = df_new[target_col]

    # Initialize a decision tree regressor model
    model = DecisionTreeRegressor()

    # Initialize RFE with the decision tree regressor model and the number of desired features (k)
    rfe = RFE(estimator=model, n_features_to_select=k)

    # Perform recursive feature elimination
    rfe.fit(df_new.drop(target_col, axis=1), y)

    # Get the selected feature indices
    selected_indices = rfe.get_support(indices=True)

    # Get the names of the selected features
    important_vars = df_new.drop(target_col, axis=1).columns[selected_indices].tolist()

    return important_vars

# Dataframe is named 'df_outlier_removed' and the target column is named 'Plant C5PlusYield'
k = 15 # Select the top 15 most important variables
important_vars_dectree = rfe_DecissionTreeRegressor(df_outlier_removed, 'Plant C5PlusYield', k)

# Print the List of important variables
print(important_vars_dectree)
```

['Feed N Plus 2A content', 'Reactor WAIT', 'H2 to HC', 'Reactor 1 Inlet Temp', 'Reactor 1 Delta T', 'Reactor 3 Delta T', 'Seperator Pressure', 'Net gas Hydrogen Purity', 'Coke on Spent Catalys t', 'Chloride Injection rate', 'Total Aromatics in feed', 'Total olefins in Feed', 'Reactor LHSV', '50% IBP', 'WABT']

```
In [8]: # Recrussive Feature Engineering using Linear Regression

def rfe_LinearRegression(df, target_col, k):
    # Create a copy of the original dataframe
    df_new = df.copy()

    # Get the target column values
    y = df_new[target_col]

    # Initialize a Linear regression model
    model = LinearRegression()

    # Initialize RFE with the linear regression model and the number of desired features (k)
    rfe = RFE(estimator=model, n_features_to_select=k)

    # Perform recursive feature elimination
    rfe.fit(df_new.drop(target_col, axis=1), y)

    # Get the selected feature indices
    selected_indices = rfe.get_support(indices=True)

    # Get the names of the selected features
    important_vars = df_new.drop(target_col, axis=1).columns[selected_indices].tolist()

    return important_vars

# Dataframe is named 'df_outlier_removed' and the target column is named 'Plant CSPlusYield'
k = 15 # Select the top 15 most important variables
important_vars_lr = rfe_LinearRegression(df_outlier_removed, 'Plant CSPlusYield', k)

# Print the List of important variables
print(important_vars_lr)

['Feed N Plus 2A content', 'Reactor WAIT', 'H2 to HC', 'Reactor 1 Delta P', 'Reactor 2 Delta P', 'Reactor 3 Delta P', 'Reactor 4 Delta P', 'Seperator Pressure', 'Coke on Spent Catalyst', 'Chloride Injection rate', 'Total Paraffins in feed', 'Total Naphthenes in feed', 'Total Aromatics in feed', 'Total olefins in Feed', 'Reactor LHSV']
```

```
In [9]: max_length = max(len(important_vars_lr), len(important_vars_dectree))

common_vars = list(set(important_vars_lr) & set(important_vars_dectree))
common_vars += [None] * (max_length - len(common_vars))

lr_only_vars = list(set(important_vars_lr) - set(important_vars_dectree))
lr_only_vars += [None] * (max_length - len(lr_only_vars))

dectree_only_vars = list(set(important_vars_dectree) - set(important_vars_lr))
dectree_only_vars += [None] * (max_length - len(dectree_only_vars))

df_rfe = pd.DataFrame({'Common': common_vars, 'Linear Regression RFE': lr_only_vars, 'Decision Tree RFE': dectree_only_vars})

df_rfe
```


Out[9]:

	Common	Linear Regression RFE	Decision Tree RFE
0	Seperator Pressure	Total Paraffins in feed	WABT
1	H2 to HC	Reactor 3 Delta P	Reactor 1 Delta T
2	Reactor LHSV	Reactor 4 Delta P	Reactor 3 Delta T
3	Reactor WAIT	Reactor 1 Delta P	Net gas Hydrogen Purity
4	Feed N Plus 2A content	Reactor 2 Delta P	Reactor 1 Inlet Temp
5	Coke on Spent Catalyst	Total Naphthenes in feed	50% IBP
6	Chloride Injection rate	None	None
7	Total Aromatics in feed	None	None
8	Total olefins in Feed	None	None
9	None	None	None
10	None	None	None
11	None	None	None
12	None	None	None
13	None	None	None
14	None	None	None

4.2 Feature selection based on Subject knowledge and VIF

```
In [10]: ## Building a Heat map first and identifying highly correlated variable
figure(figsize=(40, 24))
dataplot = sns.heatmap(df_outlier_removed.corr(), cmap="YlGnBu", annot=True)

# displaying heatmap
plt.show()
```




```
In [11]: cols_to_drop=['Reactor 1 Inlet Temp','Reactor 2 Inlet Temp', 'Reactor 3 Inlet Temp','Reactor 4 Inlet Temp',
'Reactor 1 Delta P','Reactor 2 Delta P', 'Reactor 3 Delta P', 'Reactor 4 Delta P','Net gas Hydrogen Purity',
'Total Naphthenes in feed', 'Total Aromatics in feed','Total olefins in Feed','WABT']
len(cols_to_drop)
```

Out[11]: 13

```
In [12]: df_trimmed=df_outlier_removed.drop(cols_to_drop, axis=1)
number_of_columns_dropped = df_outlier_removed.shape[1] - df_trimmed.shape[1]
number_of_columns_dropped
```

Out[12]: 13



5.0 Creating a Linear regression model and checking VIF

```
In [14]: # Splitting the dataset into X and y
X = df_trimmed.drop(['Plant CSPlusYield'], axis = 1)
y = df_trimmed[['Plant CSPlusYield']]

# Splitting data into Train and Test (80-20)
X_train, X_test,y_train, y_test = train_test_split(X,y ,random_state=104, test_size=0.2, shuffle=True)

# Scaling the X data using StandardScaler
ss = StandardScaler()
X_train = pd.DataFrame(ss.fit_transform(X_train),columns = X_train.columns)
X_test = pd.DataFrame(ss.fit_transform(X_test),columns = X_test.columns)

#Identifying total number of independent variables
print(f'Total number of independent variables are : {X_train.shape[1]}')
Total number of independent variables are : 16

In [15]: ## Applying RFE on df_trimmed to get a starting point for the model
k = 10 # Select the top 15 most important variables
important_vars_lr = rfe_LinearRegression(df_trimmed, 'Plant CSPlusYield', k)

# Print the List of important variables
print(important_vars_lr)

['Reactor WAIT', 'H2 to HC', 'Reactor 4 Delta T', 'Seperator Pressure', 'Seperator Temperature', 'Coke on Spent Catalyst', 'Chloride Injection rate', 'Reactor LHSV', 'Feed IBP', '50% IBP']
```

```
In [16]: # Concised model- First model using all columns
relevant_cols= X_train.columns
#Model 1 (Linear regression)
X_train_sm = sm.add_constant(X_train[relevant_cols])
lr = sm.OLS(y_train.values.reshape(-1,1), X_train_sm).fit()
#lr = sm.OLS(y_train, X_train_sm).fit()
lr.summary()
```

Out[16]: OLS Regression Results

Dep. Variable:	y	R-squared:	0.797			
Model:	OLS	Adj. R-squared:	0.797			
Method:	Least Squares	F-statistic:	1103.			
Date:	Wed, 24 May 2023	Prob (F-statistic):	0.00			
Time:	19:43:09	Log-Likelihood:	-3547.8			
No. Observations:	4499	AIC:	7130.			
Df Residuals:	4482	BIC:	7239.			
Df Model:	16					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	84.9682	0.008	1.07e+04	0.000	84.953	84.984
Feed N Plus 2A content	0.0071	0.025	0.285	0.775	-0.042	0.056
Reactor WAIT	-1.2812	0.025	-51.255	0.000	-1.330	-1.232
H2 to HC	-0.3396	0.020	-17.002	0.000	-0.379	-0.300
Reactor 1 Delta T	0.0233	0.016	1.487	0.137	-0.007	0.054
Reactor 2 Delta T	-0.0164	0.012	-1.391	0.164	-0.040	0.007
Reactor 3 Delta T	0.0144	0.012	1.179	0.239	-0.010	0.038
Reactor 4 Delta T	0.1396	0.014	9.838	0.000	0.112	0.167
Seperator Pressure	0.1455	0.012	12.611	0.000	0.123	0.168
Seperator Temperature	0.0933	0.010	9.606	0.000	0.074	0.112
Recycle gas purity	0.0121	0.009	1.292	0.196	-0.006	0.031
Coke on Spent Catalyst	0.4090	0.020	20.549	0.000	0.370	0.448
Chloride Injection rate	-0.2346	0.014	-16.753	0.000	-0.262	-0.207
Total Paraffins in feed	0.0220	0.026	0.836	0.403	-0.030	0.074
Reactor LHSV	-0.0658	0.016	-4.022	0.000	-0.098	-0.034
Feed IBP	-0.0369	0.013	-2.855	0.004	-0.062	-0.012
50% IBP	0.0296	0.012	2.423	0.015	0.006	0.054
Omnibus:	722.817	Durbin-Watson:	2.019			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	8947.121			
Skew:	-0.364	Prob(JB):	0.00			
Kurtosis:	9.870	Cond. No.	10.5			

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [17]: # P values are high for some variables. Checking VIF
vif = pd.DataFrame()
vif['Features'] = X_train[relevant_cols].columns
vif['VIF'] = [variance_inflation_factor(X_train[relevant_cols].values, i) for i in range(X_train[relevant_cols].shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[17]:

	Features	VIF
12	Total Paraffins in feed	10.93
0	Feed N Plus 2A content	9.89
1	Reactor WAIT	9.88
2	H2 to HC	6.31
10	Coke on Spent Catalyst	6.26
13	Reactor LHSV	4.24
3	Reactor 1 Delta T	3.88
6	Reactor 4 Delta T	3.19
11	Chloride Injection rate	3.10
14	Feed IBP	2.63
15	50% IBP	2.37
5	Reactor 3 Delta T	2.36
4	Reactor 2 Delta T	2.20
7	Seperator Pressure	2.10
8	Seperator Temperature	1.49
9	Recycle gas purity	1.40

```
In [18]: #Dropping a few variables
variables_to_be_dropped=['Reactor 1 Delta T','Reactor 2 Delta T','Reactor 3 Delta T','Reactor 4 Delta T',
'Separator Temperature'] #dropping these as Reactor delta T have high p Value and separator temp
#is highly collinear to last reactor temperature
relevant_cols = [item for item in relevant_cols if item not in variables_to_be_dropped]
```

```
In [19]: # Fitting 2nd Model
X_train_sm = sm.add_constant(X_train[relevant_cols])
lr = sm.OLS(y_train.values.reshape(-1,1), X_train_sm).fit()
#lr = sm.OLS(y_train, X_train_sm).fit()
lr.summary()
```

Out[19]: OLS Regression Results

Dep. Variable:	y	R-squared:	0.787
Model:	OLS	Adj. R-squared:	0.786
Method:	Least Squares	F-statistic:	1506.
Date:	Wed, 24 May 2023	Prob (F-statistic):	0.00
Time:	19:43:10	Log-Likelihood:	-3662.3
No. Observations:	4499	AIC:	7349.
Df Residuals:	4487	BIC:	7426.
Df Model:	11		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	84.9682	0.008	1.04e+04	0.000	84.952	84.984
Feed N Plus 2A content	-0.0635	0.024	-2.607	0.009	-0.111	-0.016
Reactor WAIT	-1.4291	0.018	-78.366	0.000	-1.465	-1.393
H2 to HC	-0.4505	0.017	-25.773	0.000	-0.485	-0.416
Seperator Pressure	0.1224	0.012	10.572	0.000	0.100	0.145
Recycle gas purity	0.0354	0.009	3.750	0.000	0.017	0.054
Coke on Spent Catalyst	0.3859	0.020	19.480	0.000	0.347	0.425
Chloride Injection rate	-0.2219	0.014	-15.770	0.000	-0.249	-0.194
Total Paraffins in feed	-0.0901	0.023	-3.993	0.000	-0.134	-0.046
Reactor LHSV	-0.0587	0.017	-3.518	0.000	-0.091	-0.026
Feed IBP	-0.0190	0.012	-1.535	0.125	-0.043	0.005
50% IBP	0.0273	0.012	2.240	0.025	0.003	0.051

Omnibus:	641.059	Durbin-Watson:	2.019
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6649.723
Skew:	-0.320	Prob(JB):	0.00
Kurtosis:	8.921	Cond. No.	8.18

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [20]: #p value for Feed IBP is high. Lets check VIF
vif = pd.DataFrame()
vif['Features'] = X_train[relevant_cols].columns
vif['VIF'] = [variance_inflation_factor(X_train[relevant_cols].values, i) for i in range(X_train[relevant_cols].shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[20]:

	Features	VIF
0	Feed N Plus 2A content	8.94
7	Total Paraffins in feed	7.66
5	Coke on Spent Catalyst	5.90
1	Reactor WAIT	5.00
2	H2 to HC	4.60
8	Reactor LHSV	4.19
6	Chloride Injection rate	2.98
9	Feed IBP	2.31
10	50% IBP	2.24
3	Seperator Pressure	2.02
4	Recycle gas purity	1.34

```
In [21]: #Dropping a few more variables
variables_to_be_dropped=['Feed IBP'] #dropping as it has high correlation with Feed composition and 50% boiling point
relevant_cols = [item for item in relevant_cols if item not in variables_to_be_dropped]
```

```
In [22]: # Fitting Final Model
X_train_sm = sm.add_constant(X_train[relevant_cols])
lr = sm.OLS(y_train.values.reshape(-1,1), X_train_sm).fit()
#lr = sm.OLS(y_train, X_train_sm).fit()
lr.summary()
```

Out[22]: OLS Regression Results

Dep. Variable:	y	R-squared:	0.787
Model:	OLS	Adj. R-squared:	0.786
Method:	Least Squares	F-statistic:	1656.
Date:	Wed, 24 May 2023	Prob (F-statistic):	0.00
Time:	19:43:10	Log-Likelihood:	-3663.5
No. Observations:	4499	AIC:	7349.
Df Residuals:	4488	BIC:	7420.
Df Model:	10		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	84.9682	0.008	1.04e+04	0.000	84.952	84.984
Feed N Plus 2A content	-0.0762	0.023	-3.324	0.001	-0.121	-0.031
Reactor WAIT	-1.4212	0.017	-81.271	0.000	-1.455	-1.387
H2 to HC	-0.4578	0.017	-27.205	0.000	-0.491	-0.425
Seperator Pressure	0.1258	0.011	11.068	0.000	0.104	0.148
Recycle gas purity	0.0358	0.009	3.795	0.000	0.017	0.054
Coke on Spent Catalyst	0.3835	0.020	19.417	0.000	0.345	0.422
Chloride Injection rate	-0.2235	0.014	-15.921	0.000	-0.251	-0.196
Total Paraffins in feed	-0.0997	0.022	-4.594	0.000	-0.142	-0.057
Reactor LHSV	-0.0600	0.017	-3.596	0.000	-0.093	-0.027
50% IBP	0.0273	0.012	2.241	0.025	0.003	0.051

Omnibus:	629.745	Durbin-Watson:	2.019
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6514.378
Skew:	-0.303	Prob(JB):	0.00
Kurtosis:	8.864	Cond. No.	7.42

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [23]: # defining a function to Test the error of predicted vs actual value
def calculate_prediction_error(X_train, y_train, lr_model):
    # Predict the target variable using the trained model
    y_train_pred = lr_model.predict(X_train)

    # Create a DataFrame for predicted values
    y_train_pred = pd.DataFrame({'Predicted': y_train_pred})

    # Reset the index of y_train and drop the current index
    y_train_act = y_train.reset_index(drop=True)

    # Concatenate actual and predicted values
    y_train_error = pd.concat([y_train_act, y_train_pred], axis=1, join="inner")

    # Calculate the error
    y_train_error['Error'] = y_train_error['Plant CSPlusYield'] - y_train_error['Predicted']

    # Calculate the R2 score
    r2_score_value = r2_score(y_train, y_train_pred)

    # Plot the distribution of error terms
    fig = plt.figure()
    sns.distplot(y_train_error['Error'], bins=20)
    fig.suptitle('Error Terms Distribution', fontsize=10)
    plt.xlabel('Error Value', fontsize=10)
    plt.grid(True)

    return y_train_error, r2_score_value
```

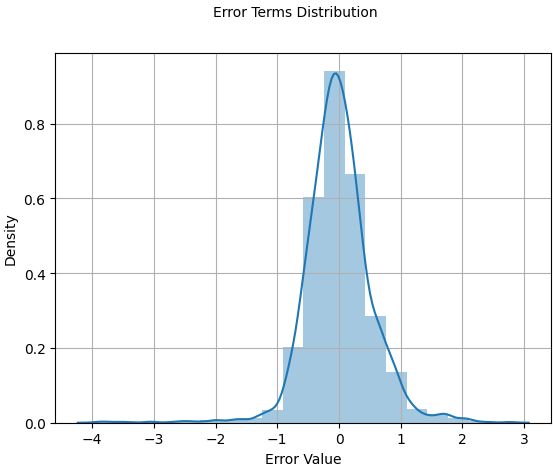


```
In [24]: # the above defined function on X_train_sm and y_train as the DataFrames, and lr is the trained Linear Regression model
y_train_error, r2_score_value = calculate_prediction_error(X_train_sm, y_train, lr)
print("R2 Score:", r2_score_value)
y_train_error.head()

R2 Score: 0.7867498508303788
```

Out[24]:

	Plant C5PlusYield	Predicted	Error
0	86.67	86.530114	0.139886
1	82.25	82.833754	-0.583754
2	82.98	83.318286	-0.338286
3	86.97	86.296869	0.673131
4	83.95	83.033286	0.916714



Checkpoint


Error terms are normally distributed.

```
In [25]: ## Fitting model on Test data and analysing
X_test_final=X_test[relevant_cols]
X_test_final.head()
X_test_sm=sm.add_constant(X_test_final)
X_test_sm.head()
# using X_test_sm and y_test as the DataFrames, and lr is the trained Linear Regression model
y_test_error, r2_score_value = calculate_prediction_error(X_test_sm, y_test, lr)
print("R2 Score:", r2_score_value)
y_test_error.head()

R2 Score: 0.7870789456093632
```

Out[25]:

	Plant C5PlusYield	Predicted	Error
0	85.72	84.497737	1.222263
1	86.27	85.877906	0.392094
2	85.94	86.052058	-0.112058
3	85.07	85.311617	-0.241617
4	84.00	83.106448	0.893552



Visualization

In [26]: *## Defining a function to plot the error analysis*


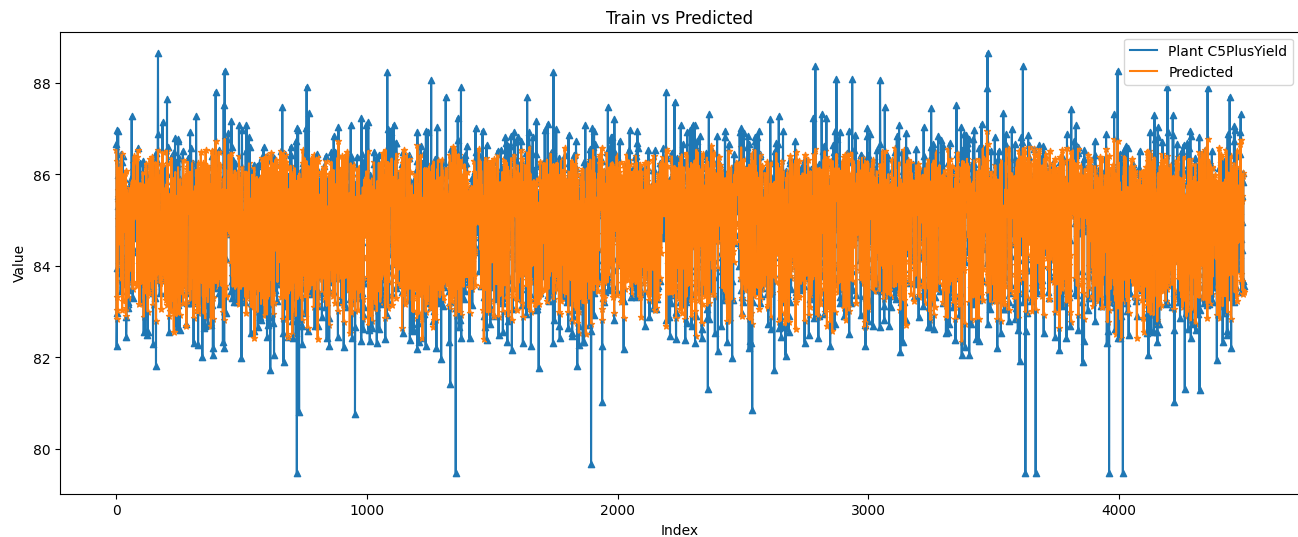
```
def plot_error_analysis(df_error,title):

    # Plot 1
    df_plot = df_error.drop('Error', axis=1)
    df_plot['x1'] = df_plot.index
    ax1 = df_plot.plot(x='x1', figsize=(16, 6))
    ax2 = df_plot.plot.scatter(x='x1', y=['Plant C5PlusYield'], marker='^', ax=ax1)
    ax3 = df_plot.plot.scatter(x='x1', y=['Predicted'], marker='*', color='ff7f0e')
    plt.title(title)
    plt.xlabel('Index')
    plt.ylabel('Value')



    # Plot 2
    plt.figure(figsize=(16, 10))
    df_error['x1'] = df_error.index
    plt.scatter(df_error['x1'], df_error['Plant C5PlusYield'], c='b', marker='^', label='Train')
    plt.scatter(df_error['x1'], df_error['Predicted'], c='r', marker='*', label='Pred')
    plt.title(title)
    plt.xlabel('Index')
    plt.ylabel('Value')
    plt.legend(loc='upper left')
    plt.show() # Add this line to show the second plot
```

In [27]: *# Plotting for Actual vs Predicted for Train data*

```
plot_error_analysis(y_train_error,'Train vs Predicted')
```



```
In [28]: # Plotting for Actual vs Predicted for Test data
plot_error_analysis(y_test_error, 'Test vs Predicted')
```



6.0 Polynomial Regression Model

```
In [29]: X_relevant=X_train_sm.drop('const',axis=1)
x=list(X_relevant.columns)
y=('Plant C5PlusYield')
x.append(y)
df_final=df_trimmed[x]
df_final.head()
```

Out[29]:

	Feed N Plus 2A content	Reactor WAIT	H2 to HC	Seperator Pressure	Recycle gas purity	Coke on Spent Catalyst	Chloride Injection rate	Total Paraffins in feed	Reactor LHSV	50% IBP	Plant C5PlusYield
0	45.46	1004.48	3.28	30.83	78.67	3.82	2.59	64.91	1.62	250.57	83.24
1	44.22	1004.54	3.17	30.32	78.67	3.56	2.64	65.89	1.62	247.78	83.20
2	44.22	1004.54	3.17	30.27	78.67	3.56	2.58	65.89	1.62	247.78	83.20
3	45.49	1002.61	3.46	32.56	77.88	3.82	2.60	63.35	1.62	249.98	83.46
4	45.46	1004.34	3.31	30.39	78.67	3.82	2.60	64.91	1.62	250.57	83.33

```
In [30]: # Create a list of dictionaries to store the results
results = []

# Test polynomial regression models from 1st to 4th order
for order in range(1, 5):
    X_train, X_test, y_train, y_test = train_test_split(df_final.drop('Plant C5PlusYield', axis=1), df_final['Plant C5PlusYield'], test_size=0.2, random_state=42)

    # Create polynomial features
    poly_features = PolynomialFeatures(degree=order)
    X_train_poly = poly_features.fit_transform(X_train)
    X_test_poly = poly_features.transform(X_test)

    # Fit the polynomial regression model
    model = LinearRegression()
    model.fit(X_train_poly, y_train)

    # Predict the target variable for train and test data
    y_train_pred = model.predict(X_train_poly)
    y_test_pred = model.predict(X_test_poly)

    # Calculate mean squared error
    mse = mean_squared_error(y_test, y_test_pred)

    # Calculate R-squared score for train and test data
    r2_train = r2_score(y_train, y_train_pred)
    r2_test = r2_score(y_test, y_test_pred)




    # Plot the actual vs. predicted values for train data
    plt.figure(figsize=(5, 3))
    plt.scatter(y_train, y_train_pred, color='blue', label='Actual vs. Predicted (Train Data)')
    plt.title(f'Polynomial Regression (Order {order})')
    plt.xlabel('Actual')
    plt.ylabel('Predicted')
    plt.legend()
    plt.grid(True)
    plt.show()

    # Plot the actual vs. predicted values for test data
    plt.figure(figsize=(5, 3))
    plt.scatter(y_test, y_test_pred, color='blue', label='Actual vs. Predicted (Test Data)')
    plt.title(f'Polynomial Regression (Order {order})')
    plt.xlabel('Actual')
    plt.ylabel('Predicted')
    plt.legend()
    plt.grid(True)
    plt.show()


    # Append the results to the list
    results.append({'Order': order, 'MSE': mse, 'R2 (Train)': r2_train, 'R2 (Test)': r2_test})

# Convert the list of dictionaries into a DataFrame
poly_r2_results = pd.DataFrame(results)


# Print the DataFrame
print(poly_r2_results)
```




Polynomial Regression (Order 2)




Polynomial Regression (Order 3)




Polynomial Regression (Order 3)



Polynomial Regression (Order 4)



Polynomial Regression (Order 4)



	Order	MSE	R2 (Train)	R2 (Test)
0	1	0.358073	0.793415	0.763524
1	2	0.278951	0.844922	0.815778
2	3	0.342419	0.898153	0.773862
3	4	0.523736	0.940705	0.654118

7.0 Ridge Regression, Lasso and Elastic Net

```
In [31]: ## Defining a function Ridge, Lasso and Elastic net
def grid_search_regression(model_name, param_grid):
    df_temp = df_final.copy()
    X = df_temp.drop('Plant C5PlusYield', axis=1)
    y = df_temp['Plant C5PlusYield']

    # Standardize the features
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)

    X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

    # Create the regressor object based on the model name
    if model_name == "Lasso":
        model = Lasso()
    elif model_name == "Ridge":
        model = Ridge()
    elif model_name == "ElasticNet":
        model = ElasticNet()
    else:
        raise ValueError("Invalid model name. Supported models: Lasso, Ridge, ElasticNet")

    # Perform grid search
    grid_search = GridSearchCV(estimator=model, param_grid=param_grid, scoring='r2', cv=10)
    grid_search.fit(X_train, y_train)

    # Train the Ridge regression model with the best hyperparameters
    best_temp = grid_search.best_estimator_
    best_temp.fit(X_train, y_train)

    # Perform cross-validation to evaluate the model
    cross_val_scores = cross_val_score(best_temp, X_scaled, y, cv=10, scoring='r2')

    # Predict on train and test data
    y_train_pred = best_temp.predict(X_train)
    y_test_pred = best_temp.predict(X_test)


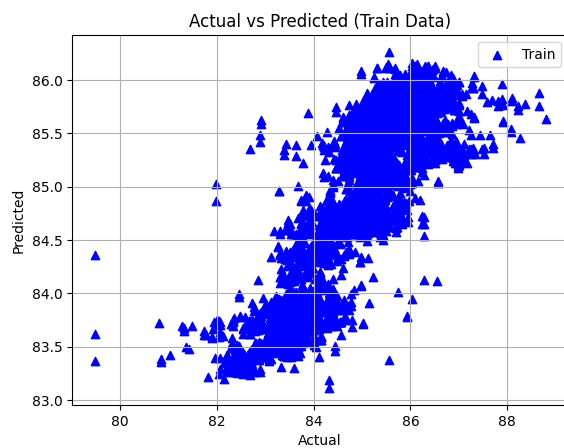
    # Calculate R2 scores for train and test data
    train_r2 = r2_score(y_train, y_train_pred)
    test_r2 = r2_score(y_test, y_test_pred)

    # Plot actual vs predicted for train data
    plt.scatter(y_train, y_train_pred, c='b', marker='^', label='Train')
    plt.xlabel("Actual")
    plt.ylabel("Predicted")
    plt.title("Actual vs Predicted (Train Data)")
    plt.legend()
    plt.grid(True)
    plt.show()

    # Plot actual vs predicted for test data
    plt.scatter(y_test, y_test_pred, c='r', marker='*', label='Test')
    plt.xlabel("Actual")
    plt.ylabel("Predicted")
    plt.title("Actual vs Predicted (Test Data)")
    plt.legend()
    plt.grid(True)
    plt.show()



    print("Best parameters: ", grid_search.best_params_)
    print("Best R-squared score: %.3f" % grid_search.best_score_)
    print("Train R2 Score: ", train_r2)
    print("Test R2 Score: ", test_r2)
    print("Cross-validation R2 Scores: ", cross_val_scores)
    print("Average Cross-validation R2 Score: ", np.mean(cross_val_scores))
    print('-'*100)
```

```
In [32]: lasso_param_grid = {'alpha': [0.1, 1.0, 10.0]}
grid_search_regression("Lasso", lasso_param_grid)
```





```
Best parameters: {'alpha': 0.1}
Best R-squared score: 0.713
Train R2 Score: 0.7151533014265671
Test R2 Score: 0.6987587245766793
Cross-validation R2 Scores: [-0.1831571  0.46838543  0.68551024  0.70536005  0.64167855  0.49062216
 0.47784368  0.49200273  0.0509346  -0.36409407]
Average Cross-validation R2 Score: 0.3465086251595856
```

```
In [33]: ridge_param_grid = {'alpha': [0.1, 1.0, 10.0]}  
grid_search_regression("Ridge", ridge_param_grid)
```



```
Best parameters: {'alpha': 1.0}  
Best R-squared score: 0.791  
Train R2 Score: 0.7934140759196404  
Test R2 Score: 0.7635365808475798  
Cross-validation R2 Scores: [0.36495254 0.53653727 0.74918257 0.77914956 0.71381132 0.64491639  
0.6168163 0.63089484 0.1917863 0.25169991]  
Average Cross-validation R2 Score: 0.5479747000196314
```

```
In [34]: elasticnet_param_grid = {'alpha': [0.1, 1.0, 10.0], 'l1_ratio': [0.25, 0.5, 0.75]}
grid_search_regression("ElasticNet", elasticnet_param_grid)
```



```
Best parameters: {'alpha': 0.1, 'l1_ratio': 0.25}
Best R-squared score: 0.748
Train R2 Score: 0.7503607356611781
Test R2 Score: 0.7291652338958736
Cross-validation R2 Scores: [ 0.02818046  0.49564942  0.73235725  0.73181379  0.6788015  0.57199081
 0.55591487  0.56718647  0.0965154 -0.13714748]
Average Cross-validation R2 Score: 0.43212624882293527
-----
```

8.0 Bagging and Boosting

8.1 XGBoost, LGBM and Bagging- Without Cross Validation

```
In [35]: def XRegressor_TrainTest(regressor):
# Loading Data
df_temp = df_final.copy()
X = df_temp.drop('Plant C5PlusYield', axis=1)
y = df_temp['Plant C5PlusYield']

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

model = regressor()

# Fit the BaggingRegressor model on the training set
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

model_name=str(regressor.__name__).upper()

# Calculate the R-squared score on the test set
test_r2 = r2_score(y_test, y_pred)*100
print(f"R-squared score on unseen data (test set) using {model_name}: %.3f" % test_r2)
print('-'*100)
```

```
In [36]: XRegressor_TrainTest(BaggingRegressor)
XRegressor_TrainTest(XGBRegressor)
XRegressor_TrainTest(LGBMRegressor)

R-squared score on unseen data (test set) using BAGGINGREGRESSOR: 90.918
-----
R-squared score on unseen data (test set) using XGBREGRESSOR: 90.351
-----
R-squared score on unseen data (test set) using LGBMREGRESSOR: 88.953
-----
```

8.2 XGBoost, LGBM and Bagging- Using K-Fold Cross Validation

```
In [37]: def XRegressor_KfoldCV(regressor):
df_temp = df_final.copy()
X = df_temp.drop('Plant_C5PlusYield', axis=1)
y = df_temp['Plant_C5PlusYield']

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Kfolds
num_folds = 10
seed = 7
kfold = KFold(n_splits=num_folds, shuffle=True, random_state=seed)

model_name=str(regressor.__name__).upper()

# Bagging Regressor
model = regressor()

results = cross_val_score(model, X_scaled, y, cv=kfold, scoring='r2')
print(f'RESULTS FOR {model_name}')
print("R-squared scores:", results)
cv_r2_mean = results.mean()
cv_r2_std = results.std()
print("Average R-squared: %.3f" % cv_r2_mean)
print("Standard Deviation of R-squared: %.3f" % cv_r2_std)

# Calculating Accuracy
cv_accuracy_mean = cv_r2_mean * 100
cv_accuracy_std = cv_r2_std * 100
print("Accuracy: %.3f%% (%.3f%%)" % (cv_accuracy_mean, cv_accuracy_std))
print('-'*30)
```

```
In [38]: XRegressor_KfoldCV(BaggingRegressor)
XRegressor_KfoldCV(XGBRegressor)
XRegressor_KfoldCV(LGBMRegressor)

RESULTS FOR BAGGINGREGRESSOR
R-squared scores: [0.93815637 0.94817044 0.89732438 0.90641677 0.92236416 0.93329892
 0.89783462 0.92795107 0.92713232 0.87621344]
Average R-squared: 0.917
Standard Deviation of R-squared: 0.021
Accuracy: 91.749% (2.113%)
-----
RESULTS FOR XGBREGRESSOR
R-squared scores: [0.93489351 0.93641588 0.86939604 0.90837882 0.92691926 0.93567633
 0.92265244 0.93162506 0.87669058 0.87511733]
Average R-squared: 0.912
Standard Deviation of R-squared: 0.026
Accuracy: 91.178% (2.615%)
-----
RESULTS FOR LGBMREGRESSOR
R-squared scores: [0.91208591 0.91885377 0.89501824 0.90680604 0.9094335 0.9179317
 0.88553086 0.91316972 0.90531064 0.87398719]
Average R-squared: 0.904
Standard Deviation of R-squared: 0.014
Accuracy: 90.381% (1.388%)
-----
```

8.3 XGBoost, LGBM and Bagging- Using Grid Search Cross Validation

```
In [39]: def grid_search_regression(model_name, param_grid):
df_temp = df_final.copy()
X = df_temp.drop('Plant C5PlusYield', axis=1)
y = df_temp['Plant C5PlusYield']

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Create the regressor object based on the model name
if model_name == "BaggingRegressor":
    model = BaggingRegressor()
elif model_name == "XGBRegressor":
    model = XGBRegressor()
elif model_name == "LGBMRegressor":
    model = LGBMRegressor()
else:
    raise ValueError("Invalid model name. Supported models: BaggingRegressor, XGBRegressor, LGBMRegressor")

# Perform grid search
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, scoring='r2', cv=10)
grid_search.fit(X_train, y_train)

# Print the best parameters and score
print("Best parameters: ", grid_search.best_params_)
print("Best R-squared score: %.3f" % grid_search.best_score_)

# Predict on the test set using the best model
best_model = grid_search.best_estimator_
y_train_pred = best_model.predict(X_train)
y_test_pred = best_model.predict(X_test)


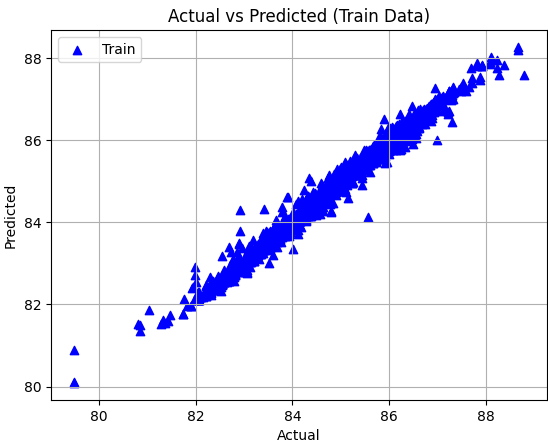
# Calculate the R-squared score on the test set
test_r2 = r2_score(y_test, y_test_pred)
print("R-squared score on unseen data (test set): %.3f" % test_r2)

# Plot actual vs predicted for train data
plt.scatter(y_train, y_train_pred, c='b', marker='^', label='Train')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted (Train Data)")
plt.legend()
plt.grid(True)
plt.show()

# Plot actual vs predicted for test data
plt.scatter(y_test, y_test_pred, c='r', marker='*', label='Test')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted (Test Data)")
plt.legend()
plt.grid(True)
plt.show()
```



```
In [40]: %%time
bagging_param_grid = {'n_estimators': [50, 100, 200], 'max_samples': [0.5, 0.8, 1.0]}
grid_search_regression("BaggingRegressor", bagging_param_grid)

Best parameters: {'max_samples': 1.0, 'n_estimators': 200}
Best R-squared score: 0.912
R-squared score on unseen data (test set): 0.918
```


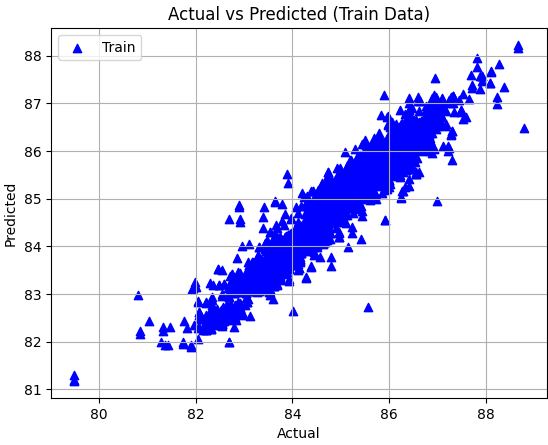


CPU times: total: 4min 11s
Wall time: 4min 12s

```
In [41]: xgb_param_grid = {'learning_rate': [0.1, 0.01, 0.001], 'max_depth': [3, 5, 7]}
grid_search_regression("XGBRegressor", xgb_param_grid)
Best parameters: {'learning_rate': 0.1, 'max_depth': 7}
Best R-squared score: 0.892
R-squared score on unseen data (test set): 0.891
```



```
In [42]: lgbm_param_grid = {'learning_rate': [0.1, 0.01, 0.001], 'max_depth': [3, 5, 7]}
grid_search_regression("LGBMRegressor", lgbm_param_grid)
Best parameters: {'learning_rate': 0.1, 'max_depth': 7}
Best R-squared score: 0.892
R-squared score on unseen data (test set): 0.885
```



END